

CSCI 4150 Project - Garbage Classification

Anushan Vimalathanan, Ken Zhou

In this project, we intend to build a model that is able to classify images of garbage into different classes. These classes are:

- cardboard
- glass
- metal
- paper
- plastic
- trash

The dataset we will be using can be found at <https://www.kaggle.com/edusdesdes/garbage-classification>

Some applications for garbage classification may be to automatically sort garbage so that garbage is placed in the correct bin. For recycling, different types of garbage may need to be recycled differently, so classifying them before recycling would help.

```
In [1]: import tensorflow as tf
import tensorflow.keras as keras
import tensorflow.keras.models as models
import tensorflow.keras.layers as layers
import tensorflow.keras.losses as losses
import tensorflow.keras.optimizers as optimizers
import cv2

import pandas as pd
import numpy as np

import matplotlib.pyplot as plt
import matplotlib.pyplot as patches
import matplotlib.image as mimg
```

Data Collection

To access the dataset in this notebook, we uploaded the dataset to google drive. In order to access it, we need to mount our drive.

```
In [2]: from google.colab import drive
drive.mount('/content/gdrive')
root_path = "/content/gdrive/My Drive/CSCI4150BU"
```

Drive already mounted at /content/gdrive; to attempt to forcibly remount, call drive.mount("/content/gdrive", force_remount=True).

The dataset contains a text file of 2 columns, which are the file name, and the class. We start by reading in the training dataset and the testing dataset into pandas DataFrames.

```
In [3]: train_path = root_path + "archive/one-indexed-files-notrash_train.txt"
test_path = root_path + "archive/one-indexed-files-notrash_test.txt"
val_path = root_path + "archive/one-indexed-files-notrash_val.txt"

# Read data into a DataFrame
train_data = pd.read_csv(train_path, delimiter=" ", header=None)
test_data = pd.read_csv(test_path, delimiter=" ", header=None)
val_data = pd.read_csv(val_path, delimiter=" ", header=None)

# Split data into x and y
x_train_text = train_data[0]
y_train = train_data[1] - 1
x_test_text = test_data[0]
y_test = test_data[1] - 1
x_val_text = val_data[0]
y_val = val_data[1] - 1
```

Since our dataset only contains the file name of the image, we need to create a function that allows us to retrieve the actual image from the file name. The function getImageFile does this by finding the path to the image from the file name, and returning a matplotlib image.

```
In [4]: # Function to get an image
def getImageFile(fileName):
    image_path = root_path + "/archive/garbage_classification/garbage_classification"

    # Add path to folder to correct image type
    if ("cardboard" in fileName):
        image_path += "/cardboard"
    elif ("glass" in fileName):
        image_path += "/glass"
    elif ("metal" in fileName):
        image_path += "/metal"
    elif ("paper" in fileName):
        image_path += "/paper"
    elif ("plastic" in fileName):
        image_path += "/plastic"
    elif ("trash" in fileName):
        image_path += "/trash"
    else:
        return "Error - File Not Found";
    image_path += "/" + fileName

    image = cv2.imread(image_path, cv2.IMREAD_COLOR)
    image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)
    #image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)
    return np.array(image)
```

We can show some images of the data by plotting the images using matplotlib

```
In [5]: # Show some images
fig, (ax1, ax2, ax3, ax4) = plt.subplots(1, 4, figsize=(20,5))

ax1.imshow(getImageFile(x_test_text[0]), cmap='gray')
ax2.imshow(getImageFile(x_test_text[1]), cmap='gray')
ax3.imshow(getImageFile(x_test_text[2]), cmap='gray')
ax4.imshow(getImageFile(x_test_text[3]), cmap='gray')

plt.show()
```

We can convert the x data for the dataset from the file name to the image by applying the getImageFile function. This will make our x data contain the image of the garbage, and our y data will contain the classification

```
In [6]: x_test = x_test_text.apply(getImageFile)
x_train = x_train_text.apply(getImageFile)
x_val = x_val_text.apply(getImageFile)

In [7]: # Stack values of the series
x_test = np.stack(x_test.values)
x_train = np.stack(x_train.values)
x_val = np.stack(x_val.values)

In [8]: #normalize the values
x_train = x_train / 255.
x_test = x_test / 255.
x_val = x_val / 255.

In [9]: x_train = tf.convert_to_tensor(x_train)
x_test = tf.convert_to_tensor(x_test)
```

We create functions that will help us visualize how our model performs.

predictTrash function will display the test images as well as the type of trash the model has predicted along with the actual type of trash in the image.

```
In [10]: def predictTrash(model):
trashTypes = ['Glass', 'Paper', 'Cardboard', 'Plastic', 'Metal', 'Trash']
prediction = model.predict(x_test)
plt.figure(figsize=(25,40))

for i in range(50):
    ax = plt.subplot(10, 5, i+1)
    plt.imshow(getImageFile(x_test_text[i]), cmap='gray')
    plt.title('Predicted: ' + trashTypes[np.argmax(prediction[i])] + " | Actual: " + trashTypes[y_test[i]])
    plt.axis('off')
```

percentDist function will display the probability distribution of an image and show the trash type of the highest probability as well as the actual trash type of the image.

```
In [11]: def percentDist(model):
prediction = model.predict(x_test)
trashTypes = ['Glass', 'Paper', 'Cardboard', 'Plastic', 'Metal', 'Trash']
plt.figure(figsize=(30,55))

for i in range(20):
    ax = plt.subplot(10, 5, i+1)
    plt.bar(trashTypes, prediction[i])
    plt.title('Predicted: ' + trashTypes[np.argmax(prediction[i])] + " | Actual: " + trashTypes[y_test[i]])
    plt.xlabel("Trash Types")
    plt.ylabel("Percentage")
```

Modelling

Now that we have the data ready, we can create a network that can be used to classify the different types of garbage.

The first model will have an input and reshape model for the data, followed by 4 16 filter convolution2D layers using a 3 by 3 kernel and 2 by 2 max pooling layer, followed by a 32 filter conv2d with another 2 by 2 max pooling layer. It will then be flattened using a layers layer and connected using two dense layers.

```
In [12]: model1 = models.Sequential([
layers.Input(shape=(384, 512)), # The images are 384x512
layers.Reshape((384, 512, 1)),
layers.Conv2D(16, (3,3), padding='same'),
layers.MaxPooling2D((2,2)), # This will make the shape 192 x 256

layers.Conv2D(16, (3,3), padding='same'),
layers.MaxPooling2D((2,2)), # This will make the shape 96 x 128

layers.Conv2D(16, (3,3), padding='same'),
layers.MaxPooling2D((2,2)), # This will make the shape 48 x 64

layers.Conv2D(16, (3,3), padding='same'),
layers.MaxPooling2D((2,2)), # This will make the shape 24 x 32

layers.Conv2D(32, (3,3), padding='same'),
layers.MaxPooling2D((2,2)), # This will make the shape 12 x 16

layers.Flatten(),
layers.Dense(64, activation='relu'),
layers.Dense(6, activation='softmax')
], name='Model_1')
```

```
In [13]: model1.compile(loss=losses.SparseCategoricalCrossentropy(), optimizer=optimizers.Adam(), metrics=['acc'])
```

```
In [14]: fitting1 = model1.fit(x_train, y_train, epochs=10, validation_data=(x_val, y_val))
```

```
Epoch 1/10
56/56 [=====] - 1865 3s/step - loss: 1.6836 - acc: 0.2729 - val_loss: 1.6413 - val_acc: 0.32
93
Epoch 2/10
56/56 [=====] - 1855 3s/step - loss: 1.4785 - acc: 0.3868 - val_loss: 1.4824 - val_acc: 0.40
95
Epoch 3/10
56/56 [=====] - 1845 3s/step - loss: 1.3398 - acc: 0.4515 - val_loss: 1.3946 - val_acc: 0.47
26
Epoch 4/10
56/56 [=====] - 1845 3s/step - loss: 1.2477 - acc: 0.5856 - val_loss: 1.2926 - val_acc: 0.52
13
Epoch 5/10
56/56 [=====] - 1845 3s/step - loss: 1.0473 - acc: 0.6184 - val_loss: 1.2977 - val_acc: 0.52
13
Epoch 6/10
56/56 [=====] - 1845 3s/step - loss: 0.8672 - acc: 0.6978 - val_loss: 1.2601 - val_acc: 0.53
96
Epoch 7/10
56/56 [=====] - 1845 3s/step - loss: 0.6924 - acc: 0.7639 - val_loss: 1.3209 - val_acc: 0.57
02
Epoch 8/10
56/56 [=====] - 1845 3s/step - loss: 0.5859 - acc: 0.7909 - val_loss: 1.5311 - val_acc: 0.55
18
Epoch 9/10
56/56 [=====] - 1845 3s/step - loss: 0.3911 - acc: 0.8615 - val_loss: 1.6103 - val_acc: 0.55
18
Epoch 10/10
56/56 [=====] - 1855 3s/step - loss: 0.2498 - acc: 0.9146 - val_loss: 1.6938 - val_acc: 0.54
57
```

```
In [15]: model1.evaluate(x_test, y_test)
```

```
14/14 [=====] - 145 076ms/step - loss: 1.9443 - acc: 0.5429
```

```
Out[15]: [1.9443359375, 0.5429234584699707]
```

Following our first model, we reduce the number of filters of the later convolution layers to see what effect it would have on the

```
In [16]: model2 = models.Sequential([
layers.Input(shape=(384, 512)), # The images are 384x512
layers.Reshape((384, 512, 1)),
layers.Conv2D(16, (3,3), padding='same'),
layers.MaxPooling2D((2,2)), # This will make the shape 192 x 256

layers.Conv2D(16, (3,3), padding='same'),
layers.MaxPooling2D((2,2)), # This will make the shape 96 x 128

layers.Conv2D(8, (3,3), padding='same'),
layers.MaxPooling2D((2,2)), # This will make the shape 48 x 64

layers.Conv2D(8, (3,3), padding='same'),
layers.MaxPooling2D((2,2)), # This will make the shape 24 x 32

layers.Conv2D(4, (3,3), padding='same'),
layers.MaxPooling2D((2,2)), # This will make the shape 12 x 16

layers.Flatten(),
layers.Dense(64, activation='relu'),
layers.Dense(6, activation='softmax')
], name='Model_2')
```

```
In [17]: model2.compile(loss=losses.SparseCategoricalCrossentropy(), optimizer=optimizers.Adam(), metrics=['acc'])
```

```
fitting2 = model2.fit(x_train, y_train, epochs=10, validation_data=(x_val, y_val))
```

```
Epoch 1/10
56/56 [=====] - 1825 3s/step - loss: 1.7553 - acc: 0.2156 - val_loss: 1.6236 - val_acc: 0.30
18
Epoch 2/10
56/56 [=====] - 1825 3s/step - loss: 1.5561 - acc: 0.3500 - val_loss: 1.5299 - val_acc: 0.39
02
Epoch 3/10
56/56 [=====] - 1825 3s/step - loss: 1.4661 - acc: 0.4120 - val_loss: 1.4301 - val_acc: 0.42
38
Epoch 4/10
56/56 [=====] - 1835 3s/step - loss: 1.3591 - acc: 0.4524 - val_loss: 1.3739 - val_acc: 0.48
48
Epoch 5/10
56/56 [=====] - 1825 3s/step - loss: 1.2620 - acc: 0.5274 - val_loss: 1.2873 - val_acc: 0.53
52
Epoch 6/10
56/56 [=====] - 1855 3s/step - loss: 1.1209 - acc: 0.5767 - val_loss: 1.2830 - val_acc: 0.52
74
Epoch 7/10
56/56 [=====] - 1825 3s/step - loss: 1.0901 - acc: 0.5893 - val_loss: 1.2909 - val_acc: 0.52
74
Epoch 8/10
56/56 [=====] - 1835 3s/step - loss: 0.9720 - acc: 0.6370 - val_loss: 1.3656 - val_acc: 0.53
35
Epoch 9/10
56/56 [=====] - 1835 3s/step - loss: 0.8665 - acc: 0.6928 - val_loss: 1.3473 - val_acc: 0.58
84
Epoch 10/10
56/56 [=====] - 1835 3s/step - loss: 0.6904 - acc: 0.7373 - val_loss: 1.3224 - val_acc: 0.57
62
```

```
In [17]: model2.evaluate(x_test, y_test)
```

```
14/14 [=====] - 145 954ms/step - loss: 1.4888 - acc: 0.5408
```

```
Out[17]: [1.488792495727539, 0.540803224627991]
```

We also want to see how increasing the number of filters the first model will affect the output of our model

```
In [18]: model3 = models.Sequential([
layers.Input(shape=(384, 512)), # The images are 384x512
layers.Reshape((384, 512, 1)),
layers.Conv2D(32, (3,3), padding='same'),
layers.MaxPooling2D((2,2)), # This will make the shape 192 x 256

layers.Conv2D(32, (3,3), padding='same'),
layers.MaxPooling2D((2,2)), # This will make the shape 96 x 128

layers.Conv2D(32, (3,3), padding='same'),
layers.MaxPooling2D((2,2)), # This will make the shape 48 x 64

layers.Conv2D(32, (3,3), padding='same'),
layers.MaxPooling2D((2,2)), # This will make the shape 24 x 32

layers.Conv2D(64, (3,3), padding='same'),
layers.MaxPooling2D((2,2)), # This will make the shape 12 x 16

layers.Flatten(),
layers.Dense(64, activation='relu'),
layers.Dense(6, activation='softmax')
], name='Model_3')
```

```
In [19]: model3.compile(loss=losses.SparseCategoricalCrossentropy(), optimizer=optimizers.Adam(), metrics=['acc'])
```

```
history3 = model3.fit(x_train, y_train, epochs=10, validation_data=(x_val, y_val))
```

```
Epoch 1/10
56/56 [=====] - 3585 6s/step - loss: 1.8914 - acc: 0.2285 - val_loss: 1.5469 - val_acc: 0.36
89
Epoch 2/10
56/56 [=====] - 3595 6s/step - loss: 1.5506 - acc: 0.3586 - val_loss: 1.5137 - val_acc: 0.34
45
Epoch 3/10
56/56 [=====] - 3555 6s/step - loss: 1.4427 - acc: 0.3872 - val_loss: 1.4625 - val_acc: 0.42
99
Epoch 4/10
56/56 [=====] - 3565 6s/step - loss: 1.3653 - acc: 0.4516 - val_loss: 1.4219 - val_acc: 0.42
99
Epoch 5/10
56/56 [=====] - 3565 6s/step - loss: 1.3126 - acc: 0.4719 - val_loss: 1.3766 - val_acc: 0.49
39
Epoch 6/10
56/56 [=====] - 3575 6s/step - loss: 1.1240 - acc: 0.5675 - val_loss: 1.3455 - val_acc: 0.46
95
Epoch 7/10
56/56 [=====] - 3615 6s/step - loss: 0.9855 - acc: 0.6303 - val_loss: 1.3312 - val_acc: 0.50
09
Epoch 8/10
56/56 [=====] - 3555 6s/step - loss: 0.7706 - acc: 0.7269 - val_loss: 1.5880 - val_acc: 0.46
65
Epoch 9/10
56/56 [=====] - 3575 6s/step - loss: 0.6510 - acc: 0.7767 - val_loss: 1.4748 - val_acc: 0.49
78
Epoch 10/10
56/56 [=====] - 3565 6s/step - loss: 0.4606 - acc: 0.8405 - val_loss: 1.7483 - val_acc: 0.53
52
```

```
In [19]: model3.evaluate(x_test, y_test)
```

```
14/14 [=====] - 225 2s/step - loss: 1.7312 - acc: 0.5197
```

```
Out[19]: [1.731215476089746, 0.5197215676307678]
```

Finally, we tried reducing the number of filters on every convolution layer in our model to see what effect it would have

```
In [20]: model4 = models.Sequential([
layers.Input(shape=(384, 512)), # The images are 384x512 with 3 colours
layers.Reshape((384, 512, 1)),
layers.Conv2D(8, (3,3), padding='same'),
layers.MaxPooling2D((2,2)), # This will make the shape 192 x 256

layers.Conv2D(8, (3,3), padding='same'),
layers.MaxPooling2D((2,2)), # This will make the shape 96 x 128

layers.Conv2D(8, (3,3), padding='same'),
layers.MaxPooling2D((2,2)), # This will make the shape 48 x 64

layers.Conv2D(8, (3,3), padding='same'),
layers.MaxPooling2D((2,2)), # This will make the shape 24 x 32

layers.Flatten(),
layers.Dense(64, activation='relu'),
layers.Dense(6, activation='softmax')
], name='Model_4')
```

```
In [21]: model4.compile(loss=losses.SparseCategoricalCrossentropy(), optimizer=optimizers.Adam(), metrics=['acc'])
```

```
In [22]: fitting4 = model4.fit(x_train, y_train, epochs=10, validation_data=(x_val, y_val))
```

```
Epoch 1/10
56/56 [=====] - 1355 2s/step - loss: 1.7400 - acc: 0.2463 - val_loss: 1.6516 - val_acc: 0.30
18
Epoch 2/10
56/56 [=====] - 1345 2s/step - loss: 1.6138 - acc: 0.3303 - val_loss: 1.5255 - val_acc: 0.39
02
Epoch 3/10
56/56 [=====] - 1345 2s/step - loss: 1.5101 - acc: 0.3791 - val_loss: 1.5956 - val_acc: 0.36
89
Epoch 4/10
56/56 [=====] - 1345 2s/step - loss: 1.4081 - acc: 0.4504 - val_loss: 1.4739 - val_acc: 0.41
77
Epoch 5/10
56/56 [=====] - 1345 2s/step - loss: 1.3488 - acc: 0.4816 - val_loss: 1.3253 - val_acc: 0.45
12
Epoch 6/10
56/56 [=====] - 1345 2s/step - loss: 1.2512 - acc: 0.5234 - val_loss: 1.2984 - val_acc: 0.46
95
Epoch 7/10
56/56 [=====] - 1345 2s/step - loss: 1.1333 - acc: 0.5728 - val_loss: 1.2651 - val_acc: 0.50
38
Epoch 8/10
56/56 [=====] - 1345 2s/step - loss: 1.0799 - acc: 0.6890 - val_loss: 1.1975 - val_acc: 0.53
96
Epoch 9/10
56/56 [=====] - 1345 2s/step - loss: 0.9399 - acc: 0.6697 - val_loss: 1.2629 - val_acc: 0.53
68
Epoch 10/10
56/56 [=====] - 1345 2s/step - loss: 0.8550 - acc: 0.7801 - val_loss: 1.2875 - val_acc: 0.56
48
```

```
In [23]: model4.evaluate(x_test, y_test)
```

```
14/14 [=====] - 125 838ms/step - loss: 1.3618 - acc: 0.5290
```

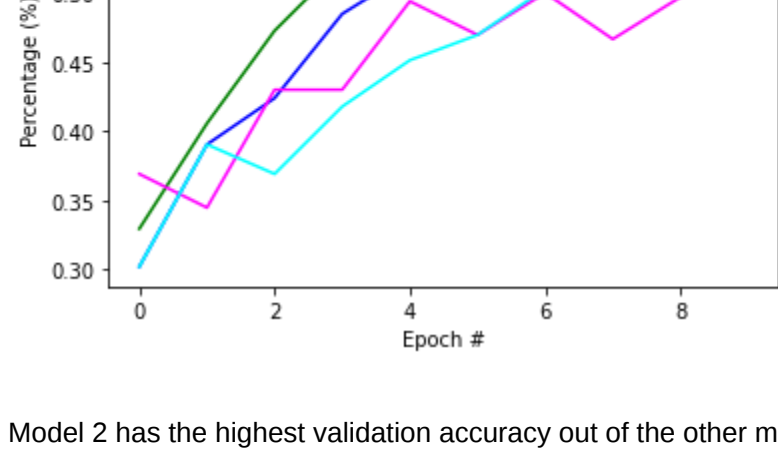
```
Out[23]: [1.361825898342896, 0.52908238684552]
```

Observations

We want to pick the model that is best suited for our study. The model we chose would be the one with the highest validation accuracy.


```
In [33]: fitList = [fitting1, fitting2, history3, fitting4]
color = ["Green", "Blue", "Magenta", "Cyan"]

for i in range(len(fitList)):
    plt.plot(fitList[i].history['val_acc'], label="Model " + str(i+1), color=color[i])
    plt.ylabel("Epoch #")
    plt.ylabel("Percentage (%)")
    plt.title("Validation Accuracy")
    plt.legend()
    plt.show()
```



Model 2 has the highest validation accuracy out of the other models.

We want to visualize how our model performs on the test dataset. The predictTrash function will show us an image of garbage with the predicted trash type and the actual trash type. The percentDist function will show us the probability distribution of the types of trash for a test and give us the predicted trash type and the actual trash type.



Conclusion

Our goal for this project was to classify images of garbage into 6 different types. After training some models, our best model was able to achieve an accuracy of 54.00% on the test data.